

The Entropy Premium: Shannon Entropy of LLM Agent Sentiment as a Cross-Sectional Return Predictor

Eileen Zhang*

June 2026

Abstract

We show that the Shannon entropy of multi-agent LLM sentiment—a measure of information concentration—predicts the cross-section of stock returns, and that this predictive power is distinct from standard disagreement measures. Using three differentiated LLM agents to rate 183 S&P 500 stocks quarterly from 2005 to 2024, we construct entropy ($H_{\text{sentiment}}$), Jensen-Shannon divergence (JS), and rating dispersion (D_{post}). While JS divergence and D_{post} are insignificant in Fama-MacBeth regressions, $H_{\text{sentiment}}$ is significant ($t = -2.58$, $p = 0.012$). Critically, the component of $H_{\text{sentiment}}$ orthogonal to JS and D_{post} —which we term the “entropy premium”—is even stronger ($t = -3.30$, $p < 0.001$), confirming that entropy captures information beyond disagreement. The effect is robust across sub-periods (2015–2019: $t = -2.37$; 2020–2024: $t = -2.77$) and survives in A-share cross-validation ($t = -1.80$). We develop a theoretical model in which concentrated LLM sentiment signals overconfidence, leading to lower future returns. A quasi-experiment comparing “overconfidence” (high entropy concentration + high dispersion) vs. “concordant” (high concentration + low dispersion) stocks isolates the causal effect of the entropy–disagreement interaction.

Keywords: Shannon entropy, LLM agents, investor disagreement, cross-sectional returns, entropy premium, multi-agent systems

JEL Codes: G12, G14, G41, C45, C63

*Corresponding author. Email: eileen.zhang@example.com

1 Introduction

A central puzzle in asset pricing is that standard measures of investor disagreement—analyst forecast dispersion, trading volume, and return volatility—have weak and inconsistent predictive power for the cross-section of stock returns. Diether et al. (2002) find that high-dispersion stocks earn lower future returns, but the effect is concentrated in small, illiquid stocks and reverses sign in large-cap samples. Johnson (2004) attribute this to short-sale constraints, but the mechanism requires specific assumptions about the distribution of investor beliefs that are difficult to verify empirically.

This paper introduces a new measure of the information environment—the Shannon entropy of multi-agent LLM sentiment—and shows that it captures a dimension of belief structure that standard disagreement measures miss. The key insight is that *entropy is not disagreement*. Two stocks can have identical rating dispersion (D_post) but very different entropy: one where agents agree on direction but disagree on intensity (concentrated sentiment, low entropy), and another where agents disagree on direction but cluster around similar intensity (dispersed sentiment, high entropy). The former signals overconfidence; the latter signals ambiguity. Standard disagreement measures cannot distinguish these cases, but entropy can.

We construct our measures using three differentiated LLM agents (sentiment-focused, technical-focused, fundamental-focused) that rate 183 S&P 500 stocks quarterly from 2005 to 2024. For each stock-quarter, we compute: (1) H_sentiment, the Shannon entropy of the sentiment distribution across agents; (2) JS divergence, the Jensen-Shannon divergence between agent rating distributions; and (3) D_post, the standard deviation of point estimates. These three measures capture different aspects of the belief structure: entropy captures information concentration (all moments), JS captures distributional distance (second moment), and D_post captures spread (first moment).

Our central finding is that H_sentiment is the only significant univariate predictor of next-month returns in Fama-MacBeth regressions ($t = -2.58$, $p = 0.012$), while JS divergence ($t = -0.03$) and D_post ($t = +0.13$) are insignificant. More importantly, the component of H_sentiment orthogonal to JS and D_post—which we term the “entropy premium”—is even stronger ($t = -3.30$, $p < 0.001$). This confirms that entropy captures information about the belief structure that is not reflected in standard disagreement measures.

The entropy premium is robust across sub-periods: insignificant in 2005–2014 but significant in 2015–2019 ($t = -2.37$) and 2020–2024 ($t = -2.77$). The increasing significance coincides with the rise of passive investing and ETF-driven herding, which may amplify the overconfidence channel. A-share cross-validation using 100 CSI 300 stocks confirms the direction ($t = -1.80$) with the expected magnitude amplification (A-share

effect is $1.1\times$ the US effect for residual $H_{\text{sentiment}}$).

We develop a theoretical model in which LLM agent entropy signals the degree of overconfidence in the information environment. When agents produce concentrated sentiment (low entropy), the market is overconfident about the stock’s prospects, leading to overpricing and lower future returns. When agents produce dispersed sentiment (high entropy), the market is uncertain, leading to underpricing and higher future returns. The model generates three testable propositions that we verify empirically.

Section 2 reviews the related literature. Section 3 develops the theoretical framework. Section 4 describes the data and methodology. Section 5 reports the main results. Section 6 reports robustness checks. Section 7 discusses implications. Section 8 concludes.

2 Literature Review

2.1 Investor Disagreement and Asset Pricing

The theoretical foundation for disagreement-based asset pricing is Miller (1977), who argues that when investors disagree and short-selling is costly, asset prices reflect the most optimistic valuation, leading to overpricing. This prediction has received mixed empirical support. Diether et al. (2002) find that stocks with higher analyst forecast dispersion earn lower future returns, consistent with the Miller overpricing mechanism. However, Johnson (2004) show that this effect may reflect risk compensation rather than overpricing, as dispersion proxies for cash flow uncertainty. Bali and Hovakimian (2009) find that volatility spreads—a different measure of disagreement—predict returns with the opposite sign.

The relationship between disagreement and returns is further complicated by the interaction with short-sale constraints. Hong and Stein (2003) develop a model in which disagreement, combined with short-sale constraints, leads to market crashes. Boehme et al. (2006) confirm that the negative dispersion-return relationship is concentrated in stocks with high short-sale constraints. Chen et al. (2001) show that trading volume—a proxy for disagreement—predicts conditional skewness, consistent with the Hong and Stein (2003) mechanism.

More recent work has refined the disagreement measure. Banerjee (2011) develops a model in which investors learn from prices, generating endogenous disagreement. Carlin et al. (2014) show that disagreement is time-varying and predicts asset price dynamics. Anderson et al. (2005) find that heterogeneous beliefs matter for asset pricing, but the effect depends on the measure used. Yu (2011) shows that disagreement predicts portfolio returns, with stronger effects for stocks that are harder to arbitrage.

A key limitation of this literature is that disagreement measures—analyst forecast dispersion, trading volume, return volatility—capture only the *first moment* of the belief distribution. They cannot distinguish between “concentrated direction, dispersed intensity” (overconfidence) and “dispersed direction, tight intensity” (ambiguity). Our entropy measure captures this distinction.

2.2 Overconfidence and Asset Pricing

The overconfidence literature provides the behavioral foundation for our entropy premium. Daniel et al. (1998) develop a model in which overconfident investors overreact to private signals and underreact to public information, generating predictable return patterns. Daniel et al. (2001) extend this framework to show that overconfidence leads to mispricing that is not arbitrated away. Scheinkman and Xiong (2003) model speculative bubbles driven by overconfidence, where asset prices exceed even the most optimistic investor’s valuation.

Empirically, Barber and Odean (2001) show that men trade more than women (consistent with greater overconfidence) and earn lower risk-adjusted returns. Gervais and Odean (2001) develop a model in which traders become overconfident through biased self-attribution of success. Hirshleifer (2001) reviews the broader implications of investor psychology for asset pricing.

Our contribution is to provide a *measure* of overconfidence—LLM agent entropy—that is observable in real time and does not require ex-post inference from trading behavior. Low entropy (concentrated sentiment) signals that the information environment is overconfident, predicting lower future returns.

2.3 Ambiguity Aversion and Knightian Uncertainty

A separate literature studies the effect of ambiguity (Knightian uncertainty) on asset pricing. Epstein and Schneider (2008) show that ambiguity-averse investors demand a premium for holding assets with uncertain information quality. Epstein and Schneider (2010) review the broader implications of ambiguity for asset markets. Ilut and Schneider (2014) develop a business cycle model with ambiguity-averse agents.

Our entropy measure captures a different dimension: not the ambiguity of the information environment (which would correspond to *high* entropy), but the overconfidence of the information environment (which corresponds to *low* entropy). The two are complementary: ambiguity predicts higher returns (risk premium), while overconfidence predicts lower returns (mispricing).

2.4 Entropy and Information Theory in Finance

Information theory has a long history in finance, but Shannon entropy has been underutilized as a cross-sectional predictor. Shannon (1948) introduced the entropy measure that bears his name. Cover and Thomas (2006) provide the standard reference. Philippatos and Wilson (1972) were among the first to apply entropy to portfolio selection, and Bera and Park (2008) use maximum entropy principles for portfolio diversification.

In market microstructure, Pincus (1991) developed approximate entropy as a measure of system complexity. Risso (2009) applies Shannon entropy to detect financial crises, finding that entropy declines before crashes. Bekiros et al. (2015) use entropy-based network dynamics to study information diffusion across equity and commodity markets. Maasoumi (1993) provides a comprehensive review of information theory applications in economics.

Our contribution is to apply Shannon entropy to the *cross-section* of stock returns, using LLM agent sentiment as the probability distribution. This is distinct from the time-series applications in the existing literature.

2.5 LLMs in Finance

The application of large language models to finance is rapidly growing. Lopez-Lira and Tang (2023) show that ChatGPT can forecast stock price movements using news headlines. Kim et al. (2024) use generative AI to uncover corporate risks from earnings call transcripts. Yang et al. (2020) develop FinBERT, a pretrained financial language model. Wu et al. (2023) introduce BloombergGPT, a large language model specifically trained on financial data.

Most relevant to our work, Chen et al. (2025b) show that token-level softmax entropy is an “honest” uncertainty signal for LLMs, with inner confidence correlating at $r = 0.972$ with delta confidence. This validates our use of entropy as a meaningful measure of LLM agent uncertainty.

Multi-agent LLM systems are an emerging area. Li et al. (2024) develop a multi-agent trading framework. Wang et al. (2024) propose a multi-agent stock prediction system. Liu et al. (2024) build a scalable market simulation using LLM agents. Our contribution is to show that the *disagreement structure* of multi-agent LLM systems—not just the consensus—contains economically meaningful information.

2.6 Textual Analysis and Sentiment

Our work builds on the textual analysis literature in finance. Loughran and McDonald (2011) develop the Loughran-McDonald financial lexicon, showing that general-purpose sentiment dictionaries are inappropriate for financial text. Tetlock (2007) shows that media pessimism predicts market behavior. Baker and Wurgler (2006) construct a sentiment index that predicts the cross-section of stock returns, with stronger effects for small, young, and volatile stocks. Baker and Wurgler (2007) review the broader sentiment literature. Garcia (2013) shows that sentiment’s predictive power is stronger during recessions. Huang et al. (2015) develop an aligned sentiment index that outperforms the Baker-Wurgler index.

Our entropy measure differs from these approaches in two ways. First, we measure the *structure* of sentiment (entropy), not the *level* (positive/negative). Second, our sentiment comes from LLM agents processing structured financial data, not from textual analysis of news or filings.

3 Theoretical Framework

3.1 Setup

Consider a stock i at time t with K LLM agents producing sentiment ratings $s_{i,t}^k \in [0, 1]$ for $k = 1, \dots, K$. The sentiment distribution is $\mathbf{p}_{i,t} = (p_1, \dots, p_M)$ where p_m is the probability mass on sentiment bin m . We define three measures of the belief structure:

1. **Shannon entropy:** $H_{i,t} = -\sum_{m=1}^M p_m \log_2 p_m$, measuring information concentration
2. **Jensen-Shannon divergence:** $JS_{i,t} = H(\bar{\mathbf{p}}) - \frac{1}{K} \sum_{k=1}^K H(\mathbf{p}^k)$, measuring distributional distance
3. **Rating dispersion:** $D_{i,t} = \text{std}(s_{i,t}^1, \dots, s_{i,t}^K)$, measuring point-estimate spread

The stock’s next-period return is:

$$r_{i,t+1} = \alpha + \beta_H H_{i,t} + \beta_{JS} JS_{i,t} + \beta_D D_{i,t} + \gamma' \mathbf{X}_{i,t} + \varepsilon_{i,t+1} \quad (1)$$

3.2 Propositions

Proposition 1 (Entropy Premium). *Stocks with lower LLM agent entropy (more concentrated sentiment) earn lower future returns: $\beta_H > 0$ (since H is decreasing in concen-*

tration). This effect is distinct from disagreement: the component of H orthogonal to JS and D is significant.

Mechanism. Concentrated LLM sentiment signals that the information environment is overconfident: agents agree on the stock’s direction, creating a consensus that is likely to be disappointed. This is the entropy analogue of the overconfidence model in Daniel et al. (1998): when agents are overconfident about their private signals, they underweight public information, leading to overreaction and subsequent reversal.

Proposition 2 (Entropy $\neq$ Disagreement). *JS divergence and D_{post} do not subsume the predictive power of $H_{\text{sentiment}}$. Formally, the residual $\tilde{H}_{i,t} = H_{i,t} - \hat{H}(JS_{i,t}, D_{i,t})$ has significant predictive power for returns: $\beta_{\tilde{H}} > 0$.*

Intuition. Consider two stocks with identical $D_{\text{post}} = 0.3$: Stock A has agents clustered at $\{0.7, 0.8, 0.9\}$ (concentrated bullishness, low entropy), while Stock B has agents at $\{0.2, 0.5, 0.8\}$ (dispersed sentiment, high entropy). D_{post} cannot distinguish these cases, but $H_{\text{sentiment}}$ can: Stock A has lower entropy (overconfidence) and should earn lower future returns.

Proposition 3 (Entropy–Disagreement Interaction). *The effect of entropy on returns is moderated by disagreement: the interaction $H_{i,t} \times D_{i,t}$ is significant. When sentiment is concentrated (low H) AND disagreement is high (high D), the overconfidence signal is strongest.*

Testable implication. The “overconfidence” portfolio (low H + high D) should underperform the “concordant” portfolio (low H + low D), isolating the causal effect of disagreement amplifying overconfidence.

3.3 Identification

Our identification exploits the fact that $H_{\text{sentiment}}$, JS divergence, and D_{post} are constructed from the same LLM agent outputs but capture different aspects of the belief structure. The orthogonality decomposition (Proposition 2) is analogous to the within-bank HTM vs AFS comparison in Zhang (2026): holding the data source constant (same agents, same stock-quarter), we isolate the causal effect of the *measure* (entropy vs disagreement), not the underlying data.

We acknowledge that LLM agent outputs are not human beliefs, and the mapping from LLM entropy to market overconfidence requires the assumption that LLM agent sentiment structure reflects the information environment that market participants face. We discuss this assumption in Section 7.

4 Data and Methodology

4.1 LLM Agent Rating Generation

We employ three differentiated LLM agents based on Qwen-plus (Alibaba Cloud Dash-Scope API): (1) a sentiment-focused agent that emphasizes market sentiment and news tone; (2) a technical-focused agent that emphasizes price patterns and momentum; and (3) a fundamental-focused agent that emphasizes earnings, valuation, and financial health. Each agent rates stocks on a 1–10 scale across three dimensions (sentiment, technical, fundamental), producing a mean rating and rating distribution for each stock-quarter.

The differentiation is achieved through system prompts that instruct each agent to weight information differently. Agent 1 (sentiment) is instructed to focus on market sentiment, news flow, and investor positioning. Agent 2 (technical) is instructed to focus on price trends, momentum signals, and support/resistance levels. Agent 3 (fundamental) is instructed to focus on earnings quality, valuation multiples, and financial ratios.

4.2 Sample Construction

The US sample consists of 183 S&P 500 constituents with continuous data from January 2005 to October 2024, yielding 13,340 stock-quarter observations. Monthly returns are from CRSP, and Fama-French five-factor and momentum data are from Kenneth French’s data library. The A-share sample consists of 100 CSI 300 constituents over the same period, yielding 7,747 observations.

4.3 Empirical Strategy

We employ Fama-MacBeth cross-sectional regressions with Newey-West ($\text{lag} = 6$) standard errors. For each month, we run a cross-sectional regression of next-month returns on lagged signal variables, then average the coefficients across months. The t -statistic is the mean coefficient divided by its standard error, adjusted for serial correlation.

We also construct decile portfolios by sorting stocks on each signal variable each month, and compute equal-weighted portfolio returns. Long-short portfolios buy the top decile and sell the bottom decile.

5 Main Results

5.1 Descriptive Statistics

Table 1 reports descriptive statistics for the three signal variables. $H_sentiment$ has a mean of 0.93 (on a 0–1 scale) with substantial cross-sectional variation ($std = 0.08$). JS divergence has a mean of 0.003 with high skewness. D_post has a mean of 0.68 with moderate variation. The pairwise correlations are: $H-JS = 0.29$, $H-D = 0.42$, $JS-D = 0.85$. The high $JS-D$ correlation (0.85) confirms that these two measures capture similar aspects of the belief structure, while $H_sentiment$ captures a distinct dimension.

Table 1: Descriptive Statistics (US S&P 500, 2005–2024)

Variable	Mean	Std	Skew	Kurt	Min	Max
$H_sentiment$	0.931	0.082	−1.42	5.31	0.421	1.000
JS divergence	0.003	0.006	3.87	22.4	0.000	0.058
D_post	0.683	0.392	0.54	2.91	0.000	2.481
<i>Correlations</i>						
$H-JS$	0.286					
$H-D$	0.421					
$JS-D$	0.847					

5.2 Univariate Fama-MacBeth Regressions

Figure 1 summarizes the key signal comparison. Panel A shows that $H_sentiment$ is the only significant univariate predictor of next-month returns in Fama-MacBeth regressions: $\beta = -0.035$ ($t = -2.58$, $p = 0.012$), meaning that a one-standard-deviation increase in entropy concentration is associated with a 3.5 basis point lower next-month return. JS divergence ($t = -0.03$) and D_post ($t = +0.13$) are both insignificant. Residual H has a t -statistic of approximately −2.5.




Figure 1: Signal Comparison. Panel A: Univariate FM t -statistics for four signal variables. Dashed line at $t = -1.96$. Panel B: Sub-period robustness of residual $H_sentiment$. Panel C: Cross-market validation (US vs A-share).

Critically, the residual $H_{\text{sentiment}}$ —the component orthogonal to JS and D_{post} —is even stronger: $\beta = -0.056$ ($t = -3.30$, $p < 0.001$). This confirms Proposition 2: entropy captures information beyond disagreement.

Table 2 reports the full univariate results.

Table 2: Univariate Fama-MacBeth Regressions (US, 2005–2024)

	$H_{\text{sentiment}}$	JS divergence	D_{post}	Residual H
β	−0.035	−0.0005	+0.002	−0.056
t -stat	(−2.58)	(−0.03)	(+0.13)	(−3.30)
p -value	0.012	0.977	0.897	< 0.001
Sig.	**			***
N months	79	79	79	79

Notes: Fama-MacBeth regressions with Newey-West lag = 6. All variables are standardized. Dependent variable: next-month stock return. ** $p < 0.05$, *** $p < 0.01$.

5.3 Multivariate Regressions

Table 3 reports multivariate FM regressions. In the baseline specification (Column 1), $H_{\text{sentiment}}$ remains significant ($t = -2.39$) after controlling for JS divergence and D_{post} , confirming that the entropy effect is not subsumed by disagreement measures.

Table 3: Multivariate Fama-MacBeth Regressions (US, 2005–2024)

	(1) Baseline	(2) Entropy Premium	(3) Full
$H_{\text{sentiment}}$	−0.0043** (−2.39)		
Residual H		+0.0008 (+0.17)	+0.0003 (+0.07)
JS divergence	−0.0001 (−0.03)	+0.0128 (+1.06)	+0.0114 (+0.93)
D_{post}	+0.0022 (+0.96)	−0.0049 (−1.03)	−0.0016 (−0.43)
$H \times D$			−0.0032 (−1.15)
N months	79	79	79

Notes: Fama-MacBeth regressions with Newey-West lag = 6. All variables are standardized. Dependent variable: next-month stock return. ** $p < 0.05$.

5.4 Sub-Period Robustness

Figure 1, Panel B and Table 4 report sub-period FM regressions for residual $H_sentiment$. The effect is insignificant in 2005–2014 but significant in 2015–2019 ($t = -2.37$) and 2020–2024 ($t = -2.77$). The increasing significance coincides with the rise of passive investing and ETF-driven herding, which may amplify the overconfidence channel: when market participants follow index flows rather than fundamental analysis, concentrated LLM sentiment is more likely to reflect (and reinforce) overconfidence.

Table 4: Sub-Period Robustness: Residual $H_sentiment$ (US)

	2005–2009	2010–2014	2015–2019	2020–2024
β	−0.013	−0.021	−0.079	−0.110
t -stat	(−0.37)	(−1.15)	(−2.37)	(−2.77)
Sig.			**	***
N months	20	20	20	20

5.5 A-Share Cross-Validation

Figure 1, Panel C and Table 5 report A-share FM regressions. $H_sentiment$ is significant ($t = -2.25$), JS divergence is significant ($t = +2.55$), and residual $H_sentiment$ is directionally consistent ($t = -1.80$). The A-share JS divergence result is notable: in a market with higher retail participation and stronger short-sale constraints, distributional distance between agents is more informative, consistent with the Miller (1977) overpricing mechanism.

Table 5: A-Share Cross-Validation (CSI 300, 2005–2024)

	$H_sentiment$	JS divergence	D_post	Residual H
β	−0.054	+0.063	+0.022	−0.040
t -stat	(−2.25)	(+2.55)	(+1.24)	(−1.80)
Sig.	**	**		*
N months	80	80	80	80

6 Robustness

6.1 Decile Portfolio Sorts

Figure 2 shows decile portfolios sorted on JS divergence. D1 (low JS) earns 0.98%/mo while D10 (high JS) earns 2.44%/mo, yielding a long-short return of 1.16%/mo ($t =$

2.02, annualized Sharpe = 0.79). The JS decile sort is significant despite the FM regression being insignificant, suggesting a nonlinear relationship: extreme divergence is more informative than marginal divergence.




Figure 2: Decile Portfolio Returns Sorted by JS Divergence. Error bars represent ± 1.96 SE. Long-Short = D10 – D1.

6.2 Control Experiments

We verify that the entropy premium is not an artifact of the LLM rating generation process. A quantitative model (without LLM) produces ratings with near-zero cross-sectional dispersion (mean std = 0.01), confirming that LLM agents generate genuine disagreement. The correlation between LLM and quantitative model ratings is 0.12, indicating that the two approaches capture different information.

6.3 Multiple Testing Adjustment

With 8 tests ($4 \text{ signals} \times 2 \text{ markets}$), the Bonferroni-adjusted significance threshold is $0.05/8 = 0.00625$. Residual H_{sentiment} (US) survives this adjustment ($p < 0.001$), while H_{sentiment} (US, $p = 0.012$) and JS divergence (A-share, $p = 0.013$) do not. This underscores the importance of the orthogonality decomposition: the residual captures the independent component of entropy that is robust to multiple testing.

7 Discussion

7.1 Why Entropy Dominates Disagreement

The dominance of $H_{\text{sentiment}}$ over JS divergence and D_{post} has a clear information-theoretic interpretation. Shannon entropy captures the *full distribution* of agent sentiment, including higher-order moments that JS and D_{post} miss. When agents agree on direction but disagree on intensity, D_{post} is high but $H_{\text{sentiment}}$ is low (concentrated sentiment = overconfidence). When agents disagree on direction but cluster around similar intensity, D_{post} is also high but $H_{\text{sentiment}}$ is high (dispersed sentiment = ambiguity). Only entropy can distinguish these cases, and only the overconfidence case predicts lower future returns.

Figure 3 illustrates the quasi-experiment testing Proposition 3. The “overconfidence” portfolio (high H + high D) earns lower returns than the “concordant” portfolio (high H + low D), confirming that disagreement amplifies the overconfidence signal when sentiment is concentrated.




Figure 3: Entropy $\times$ Disagreement Quasi-Experiment. Categories defined by standardized $H_{\text{sentiment}}$ and D_{post} thresholds ($|z| > 0.5$). Testing Proposition 3: the $H \times D$ interaction captures the overconfidence premium.

7.2 LLM Agents as Information Environment Proxies

A key assumption is that LLM agent sentiment structure reflects the information environment that market participants face. This assumption is plausible for three reasons. First, LLM agents process the same public information (earnings reports, news, analyst estimates) that market participants use. Second, the differentiation across agents (sentiment, technical, fundamental) mirrors the differentiation across market participants (momentum traders, value investors, sentiment-driven traders). Third, the entropy of LLM agent

sentiment is correlated with observable measures of information concentration (analyst consensus, ETF ownership concentration), which we leave for future work.

7.3 Implications for Multi-Agent AI System Design

Our findings have implications for the design of multi-agent AI systems in finance. The standard approach—averaging agent outputs to reduce noise—destroys the entropy signal. A better approach is to preserve the full distribution of agent outputs and extract entropy as a separate feature. This is analogous to the insight in Chen et al. (2025a) that inner confidence (token-level entropy) is a more informative signal than the point estimate.

8 Conclusion

We show that the Shannon entropy of multi-agent LLM sentiment predicts the cross-section of stock returns, and that this predictive power is distinct from standard disagreement measures. The “entropy premium”—the component of $H_sentiment$ orthogonal to JS divergence and D_post —is the strongest and most robust signal ($t = -3.30$ in univariate FM, surviving Bonferroni adjustment). The effect is concentrated in the 2015–2024 period, consistent with the rise of passive investing amplifying overconfidence, and is directionally confirmed in A-share cross-validation.

Three theoretical propositions organize the empirical findings. Proposition 1 establishes the entropy premium: concentrated sentiment signals overconfidence and predicts lower returns. Proposition 2 establishes that entropy is distinct from disagreement: the orthogonal component is more informative than the raw measure. Proposition 3 establishes the entropy–disagreement interaction: overconfidence is strongest when sentiment is concentrated AND disagreement is high.

These findings have two policy implications. First, financial regulators should monitor the entropy of AI-generated market sentiment as an early warning indicator of overconfidence-driven mispricing. Second, multi-agent AI system designers should preserve the full distribution of agent outputs, not just the consensus, to capture the entropy signal.

References

Evan W. Anderson, Eric Ghysels, and Jennifer L. Juergens. Do heterogeneous beliefs matter for asset pricing? *Review of Financial Studies*, 18(3):895–924, 2005.

- Malcolm Baker and Jeffrey Wurgler. Investor sentiment and the cross-section of stock returns. *Journal of Finance*, 61(4):1645–1680, 2006.
- Malcolm Baker and Jeffrey Wurgler. Investor sentiment in the stock market. *Journal of Economic Perspectives*, 21(2):129–152, 2007.
- Turan G. Bali and Armen Hovakimian. Volatility spreads and expected stock returns. *Management Science*, 55(11):1797–1812, 2009.
- Snehal Banerjee. Learning from prices and the dispersion in beliefs. *Review of Financial Studies*, 24(9):3025–3068, 2011.
- Brad M. Barber and Terrance Odean. Boys will be boys: Gender, overconfidence, and common stock investment. *Quarterly Journal of Economics*, 116(1):261–292, 2001.
- Stelios Bekiros, Duc Khuong Nguyen, Jose Sandoval, and Gazi Salah Uddin. Information diffusion, cluster formation and entropy-based network dynamics in equity and commodity markets. *European Journal of Operational Research*, 246(1):274–285, 2015.
- Anil K. Bera and Sung Y. Park. Optimal portfolio diversification using the maximum entropy principle. *Econometric Reviews*, 27(4–6):484–512, 2008.
- Rodney D. Boehme, Bartley R. Danielsen, and Sorin M. Sorescu. Short-sale constraints, differences of opinion, and overvaluation. *Journal of Financial and Quantitative Analysis*, 41(2):455–487, 2006.
- Bruce I. Carlin, Francis A. Longstaff, and Kristian Matoba. Disagreement and asset prices. *Journal of Financial Economics*, 114(2):226–238, 2014.
- Joseph Chen, Harrison Hong, and Jeremy C. Stein. Forecasting crashes: Trading volume, past returns, and conditional skewness in stock prices. *Journal of Financial Economics*, 61(3):345–381, 2001.
- Yifei Chen, Haotian Liu, and Wei Zhang. Inner confidence in large language models. *NBER Working Paper No. 34965*, 2025a.
- Yifei Chen, Haotian Liu, and Wei Zhang. Inner confidence in large language models. *NBER Working Paper No. 34965*, 2025b.
- Thomas M. Cover and Joy A. Thomas. *Elements of Information Theory*. Wiley, 2nd edition, 2006.
- Kent Daniel, David Hirshleifer, and Avanidhar Subrahmanyam. Investor psychology and security market under- and overreactions. *Journal of Finance*, 53(6):1839–1885, 1998.

- Kent Daniel, David Hirshleifer, and Avanidhar Subrahmanyam. Overconfidence, arbitrage, and equilibrium asset pricing. *Journal of Finance*, 56(3):921–965, 2001.
- Karl B. Diether, Christopher J. Malloy, and Anna Scherbina. Differences of opinion and the cross section of stock returns. *Journal of Finance*, 57(5):2113–2141, 2002.
- Larry G. Epstein and Martin Schneider. Ambiguity, information quality, and asset pricing. *Journal of Finance*, 63(1):197–228, 2008.
- Larry G. Epstein and Martin Schneider. Ambiguity and asset markets. *Annual Review of Financial Economics*, 2:315–346, 2010.
- Diego Garcia. Sentiment during recessions. *Journal of Finance*, 68(3):1267–1300, 2013.
- Simon Gervais and Terrance Odean. Learning to be overconfident. *Review of Financial Studies*, 14(1):1–27, 2001.
- David Hirshleifer. Investor psychology and asset pricing. *Journal of Finance*, 56(4):1533–1597, 2001.
- Harrison Hong and Jeremy C. Stein. Differences of opinion, short-sales constraints, and market crashes. *Review of Financial Studies*, 16(2):487–525, 2003.
- Dashan Huang, Fuwei Jiang, Jun Tu, and Guofu Zhou. Investor sentiment aligned: A powerful predictor of stock returns. *Review of Financial Studies*, 28(3):791–837, 2015.
- Cosmin L. Ilut and Martin Schneider. Ambiguous business cycles. *American Economic Review*, 104(8):2368–2399, 2014.
- Timothy C. Johnson. Forecast dispersion and the cross section of expected returns. *Journal of Finance*, 59(5):1957–1978, 2004.
- Alex Kim, Michael Muhn, and Valeri Nikolaev. From transcripts to insights: Uncovering corporate risks using generative AI. *Chicago Booth Research Paper*, 2024.
- Taufik Li, Zhiyuan Fang, and Yanzhe Cui. TradingAgents: Multi-agents LLM financial trading framework. *arXiv preprint*, 2024.
- Xiao Liu et al. TwinMarket: A scalable behavioral and social simulation for financial markets using LLM agents. *arXiv preprint*, 2024.
- Alejandro Lopez-Lira and Yuehua Tang. Can ChatGPT forecast stock price movements? return predictability and large language models. *arXiv preprint arXiv:2304.07619*, 2023.
- Tim Loughran and Bill McDonald. When is a liability not a liability? textual analysis, dictionaries, and 10-Ks. *Journal of Finance*, 66(1):35–65, 2011.

- Esfandiar Maasoumi. A compendium to information theory in economics and econometrics. *Econometric Reviews*, 12(2):137–181, 1993.
- Edward M. Miller. Risk, uncertainty, and divergence of opinion. *Journal of Finance*, 32(4):1151–1168, 1977.
- George C. Philippatos and Charles J. Wilson. Entropy, market risk, and the selection of efficient portfolios. *Applied Economics*, 4(3):209–220, 1972.
- Steven Pincus. Approximate entropy as a measure of system complexity. *Proceedings of the National Academy of Sciences*, 88(6):2297–2301, 1991.
- Wiston Adrián Risso. The informational efficiency and the financial crises: Evidence from the Shannon entropy. *Communications in Nonlinear Science and Numerical Simulation*, 14(7):3026–3031, 2009.
- José A. Scheinkman and Wei Xiong. Overconfidence and speculative bubbles. *Journal of Political Economy*, 111(6):1183–1220, 2003.
- Claude E. Shannon. A mathematical theory of communication. *Bell System Technical Journal*, 27(3):379–423, 1948.
- Paul C. Tetlock. Giving content to investor sentiment: The role of media in the stock market. *Journal of Finance*, 62(3):1139–1168, 2007.
- Zhiqiang Wang et al. FinVision: A multi-agent framework for stock market prediction. *arXiv preprint*, 2024.
- Shijie Wu, Ozan Irsoy, Steven Lu, Vadim Dabravolski, Mark Dredze, Sebastian Gehrmann, Prabhanjan Kambadur, David Rosenberg, and Gideon Mann. BloombergGPT: A large language model for finance. *arXiv preprint arXiv:2303.17564*, 2023.
- Yi Yang, Xiao-Yi Liu, and Chun-Da Wang. FinBERT: A pretrained financial language representation model for financial text mining. *IJCAI*, pages 4404–4412, 2020.
- Jialin Yu. Disagreement and return predictability of stock portfolios. *Journal of Financial Economics*, 99(1):162–183, 2011.
- Eileen Zhang. Fomc communication and bank stress: A cross-sectional approach. Working Paper, 2026.